

Lecture 8 - High Level Language Optimizations

DSE 512

Drew Schmidt
2022-02-17

From Last Time

- Homework is out
- Haven't made the slack channel yet
- Questions?

Basics

Profiling

- Real profiling comes later
- For now, we only need simple timing
- "Wall clock" timers
- Other kinds of timers exist...

Wallclock Timing

Python

```
import time

t0 = time.perf_counter()
time.sleep(1.2)
t1 = time.perf_counter()
t1 - t0
```

```
## 1.2235700309975073
```

R

```
system.time(Sys.sleep(1.2))[3]
```

```
## elapsed
##      1.202
```

Python Timer Class

```
import time

class Timer(object):
    def start(self):
        self.t0 = time.perf_counter()

    def stop(self):
        t1 = time.perf_counter()
        print(t1-self.t0)
```

```
t = Timer()

t.start()
time.sleep(.63)
t.stop()
```

```
## 0.6443845240864903
```

Major HLL Strategies

- Compilers and Flags
- Fundamental types
- Use efficient kernels/packages
- Vectorization
- HLL Compilers: JIT and bytecode

Compilers and Flags

- You can compile Python/R with different compilers and flags
 - `icc` instead of `gcc`
 - `-O3` instead of `-O2`
 - AND MANY MORE
- Pros:
 - Essentially **free** performance (if you can figure it out)
 - No additional effort!
- Cons:
 - Not always easy to do
 - Aggressive flags can break things
- We won't deal with this further

Fundamental Types

- Floating point: `float` (32-bit) and `double` (64-bit)
- Most computation (R, Python) automatically in double precision
- 32-bit Float:
 - half the memory
 - twice as fast (roughly)
 - not as accurate
- Python
 - Well supported in numpy
 - `np.random.rand(3).astype('f')`
 - `np.array([1, 2, 3, 4], dtype='f')`
- R
 - float package <https://cran.r-project.org/package=float>
 - fmlr package <https://hpcran.org/packages/fmlr/index.html>

Efficient Kernels

- Use existing functions/methods/packages that solve problems *well*.
 - Don't write your own linear model fitter
 - Don't write your own matrix multiplication
 - Don't write your own deep learning framework
- Exceptions
 - Interface work
 - Learning
 - Having a *genuinely* new/different approach

Vectorization

- A specific kind of use of efficient kernels
- Not rigorously defined (sometimes means something different!)
- Simplest explanation: avoiding loops
- Trades memory for runtime performance
- Operates on a vector of inputs
 - Can be a transformation: vector(s) in \rightarrow vector out (e.g. $x+1$)
 - Can be a reduction: vector(s) in \rightarrow number out (e.g. sum operator, dot product, etc)
 - Not strictly limited to these patterns...

Vectorization (Pseudocode)

Non-Vectorized

```
s = initialize()  
for (i in 1:n)  
  s += my_operation()
```

How many numbers do we store?

Vectorized

```
x = my_operation(n)  
sum(x)
```

And here?

Vectorization Example: Dot Product

$$a \cdot b = \sum_{i=1}^n a_i b_i$$

Non-Vectorized

```
d = 0
for i in range(len(a)):
    d += a[i]*b[i]
```

Vectorized

```
numpy.dot(a, b)
```

And here?

How many numbers do we store?

Bytecode

- Your computer accepts a specific kind of instructions
- Those look nothing like R/Python instructions
- Bytecode is like machine code for interpreters
- $C \xrightarrow{\text{compiler}} \text{machine code}$
- R and Python are C programs
- So R/Python $\xrightarrow{\text{compiler}} C \xrightarrow{\text{compiler}} \text{machine code}$?

Machine Code

```
#include <stdio.h>
```

```
int main()
{
    printf("hi\n");
    return 0;
}
```

```
gcc -g hello.c -o hello
```

(gdb) disass main

Dump of assembler code for function main:

```
0x00000000000001149 <+0>:    endbr64
0x0000000000000114d <+4>:    push    %rbp
0x0000000000000114e <+5>:    mov     %rsp,%rbp
0x00000000000001151 <+8>:    lea     0xeac(%rip),%rdi
0x00000000000001158 <+15>:   callq   0x1050 <puts@plt>
0x0000000000000115d <+20>:   mov     $0x0,%eax
0x00000000000001162 <+25>:   pop     %rbp
0x00000000000001163 <+26>:   retq
```

End of assembler dump.

Machine Code

```
#include <stdio.h>
```

```
int main()  
{  
    printf("hi\n");  
    return 0;  
}
```

```
gcc -g -O2 hello.c -o hello
```

(gdb) disass main

Dump of assembler code for function main:

```
0x00000000000001060 <+0>:    endbr64  
0x00000000000001064 <+4>:    sub     $0x8,%rsp  
0x00000000000001068 <+8>:    lea     0xf95(%rip),%rdi  
0x0000000000000106f <+15>:   callq   0x1050 <puts@plt>  
0x00000000000001074 <+20>:   xor     %eax,%eax  
0x00000000000001076 <+22>:   add     $0x8,%rsp  
0x0000000000000107a <+26>:   retq
```

End of assembler dump.

Python Bytecode

```
def hello():  
    print('hi')  
  
import dis  
dis.dis(hello)
```

```
##      2          0 LOAD_GLOBAL          0 (print)  
##          2 LOAD_CONST          1 ('hi')  
##          4 CALL_FUNCTION        1  
##          6 POP_TOP  
##          8 LOAD_CONST          0 (None)  
##         10 RETURN_VALUE
```

R Bytecode

```
hello = function() print("hi")
hello = compiler::cmpfun(hello)
compiler::disassemble(hello)
```

```
## list(.Code, list(12L, GETFUN.OP, 1L, PUSHCONSTARG.OP, 3L, CALL.OP,
##      0L, RETURN.OP), list(print("hi"), print, structure(c(1L,
## 9L, 1L, 30L, 9L, 30L, 1L, 1L), srcfile = <environment>, class = "srcref"),
##      "hi", structure(c(NA, 0L, 0L, 0L, 0L, 0L, 0L, 0L), class = "expressionsIndex"),
##      structure(c(NA, 2L, 2L, 2L, 2L, 2L, 2L, 2L), class = "srcrefsIndex")))
```

- Just-In-Time
- Compilation occurs during execution
- Both modern Python and R use JIT for bytecode
- Other kinds of JITs exist --- more later

Calculating Means in Python

Mean Calculation

- *Many* ways to do this
 - Stable/unstable
 - Online
 - Mean + variance
 - Online mean + variance
- Some have numerical problems that go *well* beyond the scope of this course
- Good place to start
https://en.wikipedia.org/wiki/Kahan_summation_algorithm
- Real example usage <https://github.com/wrathematics/proginfo>

Generate Some Random Data

```
import random
random.seed(1234)
import numpy as np
n = 100000
x = np.random.rand(n)
```

Naive Mean

```
def mean(x):  
    s = 0.0  
    n = len(x)  
    for i in range(n):  
        s += x[i]  
    return s / n
```

```
t.start()  
mean_x = mean(x)  
t.stop()
```

```
## 0.040856459992937744
```

Vectorized Mean

```
def mean_numpy(x):  
    return np.sum(x) / len(x)
```

```
t.start()  
mean_x = mean_numpy(x)  
t.stop()
```

```
## 0.009058568044565618
```


- A different kind of JIT compiler
- One of the cooler Python features!
- Converts some Python into machine code using LLVM
 - some core operations
 - NumPy
 - potentially good for loops of scalar/vector/matrix ops
- Somewhat similar to armacmp

<https://github.com/dirkschumacher/armacmp>



NUMBA Mean

```
from numba import jit

@jit(nopython = True)
def mean_numba(x):
    s = 0.0
    n = len(x)
    for i in range(n):
        s += x[i]
    return s / n
```

```
t.start()
mean_x = mean_numba(x)
t.stop()
```

```
## 0.2743876969907433
```

```
t.start()
mean_x = mean_numba(x)
t.stop()
```

```
## 0.007991217076778412
```


Calculating π

Monte Carlo π Simulation

$$\text{Area} = \pi r^2$$

```
set.seed(1234)
n = 1000
x = runif(n)
y = runif(n)

circ_pts = seq(0, pi/2, length.out=100)
library(ggplot2)
g = ggplot(data.frame(x=x, y=y), aes(x, y)) +
  theme_bw() +
  geom_point() +
  annotate("path",
    x = 0 + 1*cos(circ_pts),
    y = 0 + 1*sin(circ_pts),
    color="red"
  )
```



Estimating π : Naive Implementation

```
pi_sim = function(n=1e6, seed=1234)
{
  set.seed(seed)
  x = runif(n)
  y = runif(n)

  s = 0
  for (i in 1:n)
  {
    if (x[i]*x[i] + y[i]*y[i] < 1)
      s = s + 1
  }

  4 * s / n
}
```

```
system.time({
  pi_est <- pi_sim()
})[3]
```

```
## elapsed
##    0.236
```

```
pi_est
```

```
## [1] 3.143128
```

```
pi_sim_cmp <- compiler::cmpfun(pi_sim)
system.time({
  pi_est <- pi_sim_cmp()
})[3]
```

```
## elapsed
##    0.238
```

Estimating π : Vectorized

```
pi_sim_vec = function(n=1e6, seed=1234)
{
  set.seed(seed)
  x = runif(n)
  y = runif(n)
  4 * sum(x*x + y*y < 1) / n
}
```

```
system.time({
  pi_est <- pi_sim_vec()
})[3]
```

```
## elapsed
##    0.045
```

```
pi_est
```

```
## [1] 3.143128
```

Estimating π : Python

- Discussed last year →
- Ungraded Homework:
 - Create naive impl
 - Create numpy impl
 - Use numba impl
 - Time them all
 - Experiment with different sized samples

```
import random
from numba import jit

@jit(nopython = True)
def monte_carlo_pi(nsamples):
    acc = 0
    for i in range(nsamples):
        x = random.random()
        y = random.random()
        if (x**2 + y**2) < 1:
            acc += 1
    return 4 * acc / nsamples
```

Questions?