

Lecture 9 - Computational Linear Algebra Part 1

DSE 512

Drew Schmidt
2022-02-22


From Last Time

- Homework is out --- due Saturday
- Haven't made the slack channel yet
- Questions?

Linear Algebra

Linear Algebra

- LA dominates scientific and data computing
- Our focus will be on LA for data
 - All matrices are real-valued
 - Usually non-square



Linear Algebra On Your Computer

What happens when you multiply two matrices?

Python

```
np.dot(A, B)
```

R

```
A %*% B
```

Recall: Terminology

- **gemm** - matrix-matrix multiply
- **BLAS** - Basic Linear Algebra Subprograms; matrix library
- **FLOPS** - Floating Point Operations Per Second (adds and multiplies)
- **LINPACK** - Solve $Ax = b$
- **TOP500** - list of computers ranked by LINPACK benchmark



NAME

dgemm - perform one of the matrix-matrix operations $C := \alpha * \text{op}(A) * \text{op}(B) + \beta * C$

SYNOPSIS

```
SUBROUTINE DGEMM(TRANSA, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB,  
                 BETA, C, LDC)
```

```
CHARACTER * 1 TRANSA, TRANSB  
INTEGER M, N, K, LDA, LDB, LDC  
DOUBLE PRECISION ALPHA, BETA  
DOUBLE PRECISION A(LDA,*), B(LDB,*), C(LDC,*)
```

```
SUBROUTINE DGEMM_64(TRANSA, TRANSB, M, N, K, ALPHA, A, LDA, B, LDB,  
                   BETA, C, LDC)
```

```
CHARACTER * 1 TRANSA, TRANSB  
INTEGER*8 M, N, K, LDA, LDB, LDC  
DOUBLE PRECISION ALPHA, BETA  
DOUBLE PRECISION A(LDA,*), B(LDB,*), C(LDC,*)
```

Basic Linear Algebra Subprograms

- Fortran 77 library
- Basic arithmetic operations
- Level 1: vector/vector operations
- Level 2: matrix/vector operations
- Level 3: matrix/matrix operations

Linear Algebra PACKage

- Fortran 77 library
- Matrix factorizations

BLAS Library: Python

```
import numpy as np
np.show_config()
```

```
## blas_mkl_info:
##   NOT AVAILABLE
## blis_info:
##   NOT AVAILABLE
## openblas_info:
##   NOT AVAILABLE
## atlas_3_10_blas_threads_info:
##   NOT AVAILABLE
## atlas_3_10_blas_info:
##   NOT AVAILABLE
## atlas_blas_threads_info:
##   NOT AVAILABLE
## atlas_blas_info:
##   NOT AVAILABLE
## accelerate_info:
##   NOT AVAILABLE
## blas_info:
```

BLAS Library: Python

```
np.__config__.blas_info
```

```
## {'libraries': ['blas', 'blas'], 'library_dirs': ['/usr/lib/x86_64-linux-gnu'], 'include_dirs': ['
```

BLAS Library: R

```
sessionInfo()
```

```
## R version 4.1.2 (2021-11-01)
## Platform: x86_64-pc-linux-gnu (64-bit)
## Running under: Ubuntu 20.04.3 LTS
##
## Matrix products: default
## BLAS:   /usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3
## LAPACK: /usr/lib/x86_64-linux-gnu/openblas-pthread/liblapack.so.3
##
## locale:
##  [1] LC_CTYPE=en_US.UTF-8      LC_NUMERIC=C
##  [3] LC_TIME=en_US.UTF-8      LC_COLLATE=en_US.UTF-8
##  [5] LC_MONETARY=en_US.UTF-8  LC_MESSAGES=en_US.UTF-8
##  [7] LC_PAPER=en_US.UTF-8     LC_NAME=C
##  [9] LC_ADDRESS=C             LC_TELEPHONE=C
## [11] LC_MEASUREMENT=en_US.UTF-8 LC_IDENTIFICATION=C
##
## attached base packages:
## [1] stats      graphics  grDevices  datasets  utils      methods   base
```

BLAS Library: R

```
sessionInfo()$BLAS
```

```
## [1] "/usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3"
```

```
sessionInfo()$LAPACK
```

```
## [1] "/usr/lib/x86_64-linux-gnu/openblas-pthread/liblapack.so.3"
```

BLAS Interfaces

- HLL
 - Matlab
 - NumPy
 - base R
 - Basically every language
 - Many extensions ...
- C
 - lapacke
 - GSL
- C++
 - Armadillo
 - Eigen
 - Boost
 - fml

BLAS Implementations

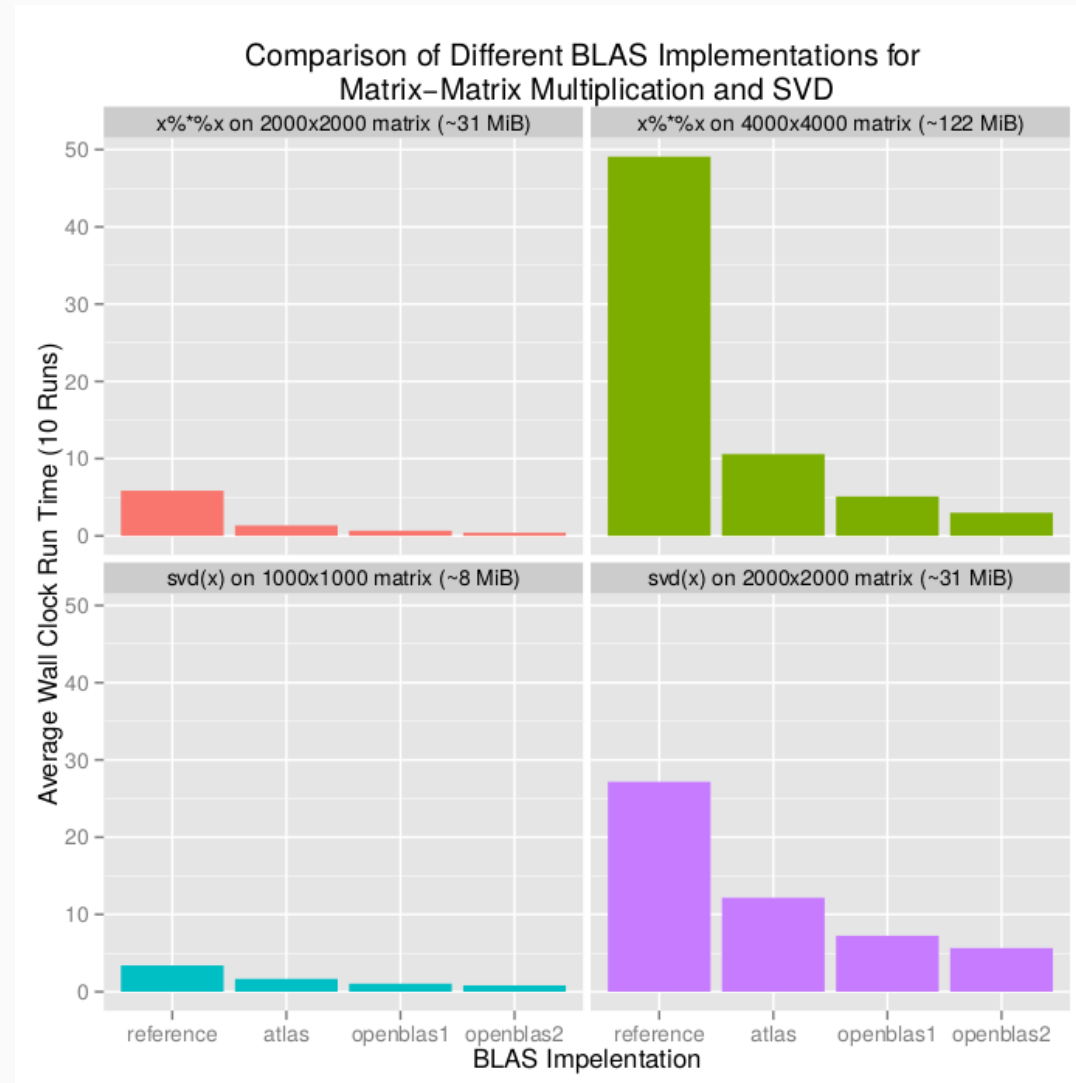
- Open source BLAS/LAPACK implementations
 - netlib (reference)
 - Atlas
 - OpenBLAS
- Proprietary BLAS/LAPACK implementations
 - Intel MKL
 - Apple Accelerate
 - AMD Optimizing CPU Libraries (AOCL)
- GPU semi-implementations
 - NVIDIA cuBLAS and NVBLAS
 - AMD roc/hip

Question

What makes a BLAS implementation "high performance"?



Some Benchmarks



The LINPACK Benchmark

- Solve the system $Ax = b$
 - A - $n \times n$ matrix (you choose n)
 - Double precision
 - Must use LU with partial pivoting
 - $A = LU$
 - $b = Ax = LUx$
 - Double triangular system solver
 - Solution must satisfy some accuracy conditions.
- $\frac{2}{3}n^3 + 2n^2$ operations
- Most FLOPS wins!

Basic LINPACK Benchmarking

- Pick an n
- Create random $A_{n \times n}$ and x
- Define $b = Ax$
- Forget x momentarily
- Get wallclock time t for solving $A\hat{x} = b$ (following the rules)
- $GFLOPS = \frac{\frac{2}{3}n^3 + 2n^2}{t \cdot 10^9}$
- Try to find the best such n

Running LINPACK

- HPL <http://www.netlib.org/benchmark/hpl/>
- Python
 - `numpy.linalg.solve(A, b)`
- R
 - `solve(A, b)`
 - [okcpuid package](#)

```
okcpuid::linpack()
```

- Your iPhone?! <https://apps.apple.com/us/app/linpack/id380883195>

My Machine

```
okcpuid::cpu_clock()
```

```
## $ncores  
## [1] 16  
##  
## $clock.os  
## 1993 MHz  
##  
## $clock.tested  
## 3393 MHz  
##  
## $peak  
## 217.152 GFLOPS
```

My LINPACK Results

- OpenBLAS with 8 threads
- N.max: 50,000
- R.max: 186.655 GFLOPS
- R.peak: 217.152 GFLOPS
- max-to-peak: 85.956%

Numerics in Regression

Linear Regression

$$\begin{aligned} y = X\beta &\iff X^T y = X^T X \beta \\ &\iff (X^T X)^{-1} X^T y = \beta \end{aligned}$$

...right?

Linear Regression

```
x = matrix(1:30, nrow=10)
y = runif(10)

solve(t(x) %*% x) %*% t(x) %*% y
```

```
## Error in solve.default(t(x) %*% x): Lapack routine dgesv: system is exactly singular: U[3,3] = 0
```

```
lm.fit(x, y)$coefficients
```

```
##           x1           x2           x3
## -0.07202751  0.06781070          NA
```


Linear Regression

```
set.seed(1234)
m = 10000
n = 250
x = matrix(rnorm(m*n), nrow=m, ncol=n)
y = runif(m)

system.time(solve(t(x) %*% x) %*% t(x) %*% y)
```

```
##      user  system elapsed
##    0.900    1.907    0.323
```

```
system.time(lm.fit(x, y))
```

```
##      user  system elapsed
##    0.348    0.003    0.352
```

Condition Numbers

[T]he condition number [...] measures how much the output value [...] can change for a small change in the input argument

In linear regression the condition number of the moment matrix can be used as a diagnostic for multicollinearity.

Source: https://en.wikipedia.org/wiki/Condition_number

What's Going On?

```
set.seed(1234)
x = matrix(rnorm(30), 10)

kappa(x)
```

```
## [1] 1.649968
```

```
kappa(x)^2
```

```
## [1] 2.722393
```

```
kappa(t(x) %*% x)
```

```
## [1] 2.891373
```

Conclusion

There are *good ways* and *bad ways* to compute the same thing.

Quick Review of Matrix Factorizations

Quick Background

- Matrix products are associative but not (in general) commutative
- rank --- number of linearly independent columns
- $(AB)^T = B^T A^T$
- $(AB)^{-1} = B^{-1} A^{-1}$
- $\det(AB) = \det(A) \det(B)$
- $A_{n \times n}$ is orthogonal if $A^T A = A A^T = I_{n \times n}$

$$A_{n \times n} = L_{n \times n} U_{n \times n}$$

- solving a square system
- (general) matrix inversion
- calculating determinants
- $A = PLU$

Cholesky

$$A_{n \times n} = L_{n \times n} L_{n \times n}^T = U_{n \times n}^T U_{n \times n}$$

- specialized LU for "positive definite" matrices
- inversion of correlation/covariance matrices

$$A_{m \times n} = Q_{m \times n} R_{n \times n}$$

- Q orthogonal ($Q^T Q = I$), R upper triangular
- solving over/under-determined systems
- linear regression
- useful in distributed contexts for PCA
- $A = LQ$

Eigendecomposition

$$A_{n \times n} = V_{n \times n} \Delta_{n \times n} V_{n \times n}^T$$

- V is orthogonal, Δ is diagonal
- Very useful to engineers; not that useful for data science
- Can be used for PCA
- Important connections to ***the other*** factorization...

Singular Value Decomposition

Singular Value Decomposition (SVD)

- ***THE*** most important matrix factorization for data
- SVD is to data what Eigendecomposition is to engineers
- SVD can do *almost anything*
 - Not always the fastest
 - Often the most accurate

"Compact" SVD

$$A_{m \times n} = U_{m \times k} \Sigma_{k \times k} V_{n \times k}^T$$

- k is usually minimum of m and n
- may be taken to be the rank of A which is no greater than the minimum of m and n
- Σ is diagonal
- U and V are orthogonal
 - $U^T U = I_{n \times n}$
 - $V^T V = I_{k \times k}$

Orthogonality of U and V

```
X = matrix(1:30, 10)
s = svd(X, nu=2, nv=2)
U = s$u
V = s$v
crossprod(U) |> round(digits=8)
```

```
##      [,1] [,2]
## [1,]    1    0
## [2,]    0    1
```

```
crossprod(V) |> round(digits=8)
```

```
##      [,1] [,2]
## [1,]    1    0
## [2,]    0    1
```

```
tcrossprod(V) |> round(digits=8)
```

```
##      [,1]      [,2]      [,3]
## [1,] 0.8333333 0.3333333 -0.1666667
## [2,] 0.3333333 0.3333333 0.3333333
## [3,] -0.1666667 0.3333333 0.8333333
```

"Full" SVD

$$A_{m \times n} = U_{m \times m} \Sigma_{m \times n} V_{n \times n}^T$$

- Rarely done in software
- $U^T U = U U^T = I_{m \times m}$
- $V^T V = V V^T = I_{n \times n}$

Connection to Eigendecomposition

$$\begin{aligned} A^T A &= (U \Sigma V^T)^T (U \Sigma V^T) \\ &= V \Sigma U^T U \Sigma V^T \\ &= V \Sigma^2 V^T \end{aligned}$$

- Likewise $AA^T = U \Sigma^2 U^T$
- We will use this extensively in the parallelism module!

Question

So what can SVD do?



SVD: Matrix Inversion

$$\begin{aligned} A^{-1} &= (U\Sigma V^T)^{-1} \\ &= V\Sigma^{-1}U^T \end{aligned}$$

SVD: System Solving

$$\begin{aligned} Ax = b &\iff U\Sigma V^T x = b \\ &\iff x = V\Sigma^{-1}U^T b \end{aligned}$$

SVD: Determinants

$$\begin{aligned} |\det(A)| &= |\det(U\Sigma V^T)| \\ &= \det(\Sigma) \\ &= \prod_{i=1}^n \sigma_{ii} \end{aligned}$$

```
x = matrix(rnorm(25), nrow=5)
det(x)
```

```
## [1] -0.9838918
```

```
prod(svd(x, nu=0, nv=0)$d)
```

```
## [1] 0.9838918
```

SVD: Regression (over/under-determined systems)

$$\begin{aligned}y = X\beta &\iff y = U\Sigma V^T\beta \\ &\iff V\Sigma^{-1}U^Ty = \beta\end{aligned}$$

SVD: Column Rank

$$\begin{aligned}\text{rank}(A) &= \text{rank}(U\Sigma V^T) \\ &= |\{\sigma \mid \sigma > 0\}| \end{aligned}$$

SVD: Condition Number

$$\text{cond}(A) = \frac{\sigma_1}{\sigma_n}$$

Ungraded Homework

- (Orthogonal matrices)
 1. Prove that a product of 2 orthogonal matrices is again orthogonal.
 2. Prove that if Q is a square orthogonal matrix, its determinant must be 1 or -1 (hint: start with $Q^T Q = I$).
 3. If a square matrix has determinant 1, is it necessarily orthogonal?
- Let A be a square, symmetric matrix of order n , and let $A = U\Sigma V^T$ be its SVD. If each singular value $\sigma_i \geq 0$, find the matrix square root of A . That is, find the matrix B so that $BB = A$.

Ungraded Homework

- Let A be a square matrix of order n . The polar decomposition is a matrix factorization where $A = UP$, where U is orthogonal and P is positive semi-definite, each of order n . Derive U and P via the SVD (hint: U_{polar} is not U_{svd}).
- Let A be a square symmetric matrix of order n .
 1. Use its eigendecomposition to calculate $\det(A)$.
 2. Using only its eigendecomposition, show that the trace of A is equal to the sum of the eigenvalues (hint: first show that by combining the definitions of trace and matrix product, $\text{tr}(AB) = \text{tr}(BA)$; or use your fancy inner product math if you know that).

Wrapup

- Next time:
 - The linear algebra of PCA and regression
 - More than you ever wanted to know!
- Resources
 - Golub, G.H. and Van Loan, C.F., 2013. Matrix computations. JHU press.
 - Rencher, A.C. and Schimek, M.G., 1997. Methods of multivariate analysis. Computational Statistics, 12(4), pp.422-422.
 - Matrix Factorizations for Data Analysis <https://fml-fam.github.io/blog/2020/07/03/matrix-factorizations-for-data-analysis/>
 - SVD Via Lanczos Iteration <https://fml-fam.github.io/blog/2020/06/15/svd-via-lanczos-iteration/>

Questions?