

Lecture 22 - Profiling Basics

DSE 512

Drew Schmidt
2022-04-19

From Last Time

- Homework 3
 - Graded
 - Let's talk about it
- Homework 4
 - Assigned
 - Let's talk about it

Where We've Been

Module 1: Basic Cloud and HPC

- Lecture 1 - Introduction
- Lecture 2 - Overview of HPC and the Cloud
- Lecture 3 - Introduction to Remote Computing
- Lecture 4 - Introduction to Containers
- Lecture 5 - Introduction to ISAAC
- Lecture 6 - MPI and Singularity

Module 2: Performance Optimization

- Lecture 7 - Introduction to Performance Optimization
- Lecture 8 - High Level Language Optimizations
- Lecture 9 - Computational Linear Algebra Part 1
- Lecture 10 - Computational Linear Algebra Part 1
- Lecture 11 - GPGPU (The Easy Parts) Part 1
- Lecture 12 - GPGPU (The Easy Parts) Part 2
- Lecture 13 - Utilizing Compiled Code
- Lecture 14 - I/O

Module 3: Parallelism

- Lecture 15 - Introduction to Parallelism
- Lecture 16 - Forks and Threads Part 1
- Lecture 17 - Forks and Threads Part 2
- Lecture 18 - MPI Part 1
- Lecture 19 - MPI Part 2
- Lecture 20 - MPI Part 3
- Lecture 21 - MapReduce

Where We're Headed

Module 4: Profiling

- Lecture 22 - Profiling Basics
- Lecture 23 - HLL Profiling
- Lecture 24 - Advanced Profiling (Hardware and MPI)


Module 5: Deep Learning

- Lecture 25 - Basic Intro
- Lecture 26 - DL for Practitioners
- Lecture 27 - Distributed Training

Profiling Basics

Profiling

- Gathering information
- Can do "profiling" on lots of different things
 - Customer profiling
 - Profiling in policing
 - ...
- We are interested in *performance profiling*



How Does Profiling Work?

Performance Profiling Examples

- Software profiling
 - a simple timer
 - line profiling
- API profiling
 - MPI operations
 - Tensorflow profiler
- Hardware profiling
 - Memory profiler
 - Hardware counters
 - CUDA



What Can Profiling Do?

FIXING the ENTIRE SM64 Source Code (INSANE N64 performance)



https://www.youtube.com/watch?v=t_rzYnXEQIE

Language Timers

- R
 - `system.time()`
 - `Sys.time()`
- Python's many time utilities
 - `time.perf_counter()`
 - ...

time

```
time Rscript -e "1+1"
```

```
## [1] 2
```

```
##
```

```
## real    0m0.287s
```

```
## user    0m0.484s
```

```
## sys     0m0.838s
```

```
time Rscript -e "x = runif(1e8)"
```

```
##
```

```
## real    0m2.203s
```

```
## user    0m2.198s
```

```
## sys     0m0.991s
```

Project home <https://github.com/wrathematics/proginfo>

- Uses sampling
- Reports basic CPU and GPU info
- MIN/MEAN/MAX (SD) / TOTAL

```
./proginfo Rscript -e '1+1'
```

```
[1] 2
```

```
## Program Info (from 27 polls)
```

```
CPU
```

```
- Wall time: 0.280  
- Utilization: 3.184%  
- RAM: 19.109/19.142/19.167 (62.791) / 62.810 GiB
```

```
GPU (CUDA=11.4 Driver=470.103.01)
```

```
- Utilization  
  + (Device 0) 6/8.96/11 (2.50) / 100%  
- RAM:  
  + (Device 0) 0.579/0.579/0.579 (0.000) / 7.926
```

Memory example

```
./proginfo Rscript -e 'x=runif\$(1e9\)'
```

```
## Program Info (from 1890 polls)
```

CPU

- Wall time: 19.454
- Utilization: 2.327%
- RAM: 19.147/22.921/26.671 (60.604) / 62.810 GiB

GPU (CUDA=11.4 Driver=470.103.01)

- Utilization
 - + (Device 0) 0/0.50/1 (0.50) / 100%
- RAM:
 - + (Device 0) 0.554/0.554/0.554 (0.000) / 7.926 GiB

GPU example

```
./proginfo Rscript -e 'suppressMessages(library(fmlr)); c = card(); x = gpumat(c, 25000,
```

```
# gpumat 25000x25000 type=f
```

```
## Program Info (from 381 polls)
```

CPU

- Wall time: 3.949
- Utilization: 2.606%
- RAM: 19.122/19.359/19.509 (62.692) / 62.810 GiB

GPU (CUDA=11.4 Driver=470.103.01)

- Utilization
 - + (Device 0) 0/1.94/27 (4.76) / 100%
- RAM:
 - + (Device 0) 0.554/0.673/3.105 (0.189) / 7.926 GiB

Basic Memory Profiling

- Can be quite sophisticated
 - valgrind
 - gdb
- Or also very simple

```
library(memuse)  
Sys.procmem()
```

```
## Size: 95.207 MiB  
## Peak: 95.207 MiB
```

```
x = runif(1e8)  
mu(x)
```

```
## 762.939 MiB
```

```
rm(x);invisible(gc())  
Sys.procmem()
```

```
## Size: 95.477 MiB  
## Peak: 858.188 MiB
```

Line Profiling

```
m = 10000
n = 250
x = matrix(rnorm(
Rprof()
pca = prcomp(x)
Rprof(NULL)
summaryRprof()
```

\$by.self

	self.time	self.pct	total.time	total.pct
"La.svd"	0.68	69.39	0.72	73.47
"%*%"	0.12	12.24	0.12	12.24
"aperm.default"	0.04	4.08	0.04	4.08
"array"	0.04	4.08	0.04	4.08
"matrix"	0.04	4.08	0.04	4.08
"sweep"	0.02	2.04	0.10	10.20

output truncated by presenter

\$by.total

	total.time	total.pct	self.time	self.pct
"prcomp"	0.98	100.00	0.00	0.00
"prcomp.default"	0.98	100.00	0.00	0.00
"svd"	0.76	77.55	0.00	0.00
"La.svd"	0.72	73.47	0.68	69.39

output truncated by presenter

\$sample.interval

[1] 0.02

Line Profiling

```
m = 10000
n = 250
x = matrix(rnorm(

Rprof(interval=.9
pca = prcomp(x)
Rprof(NULL)

summaryRprof()
```

```
$by.self
```

```
[1] self.time self.pct total.time total.pct
<0 rows> (or 0-length row.names)
```

```
$by.total
```

```
[1] total.time total.pct self.time self.pct
<0 rows> (or 0-length row.names)
```

```
$sample.interval
```

```
[1] 0.99
```

```
$sampling.time
```

```
[1] 0
```

Profiling Overhead

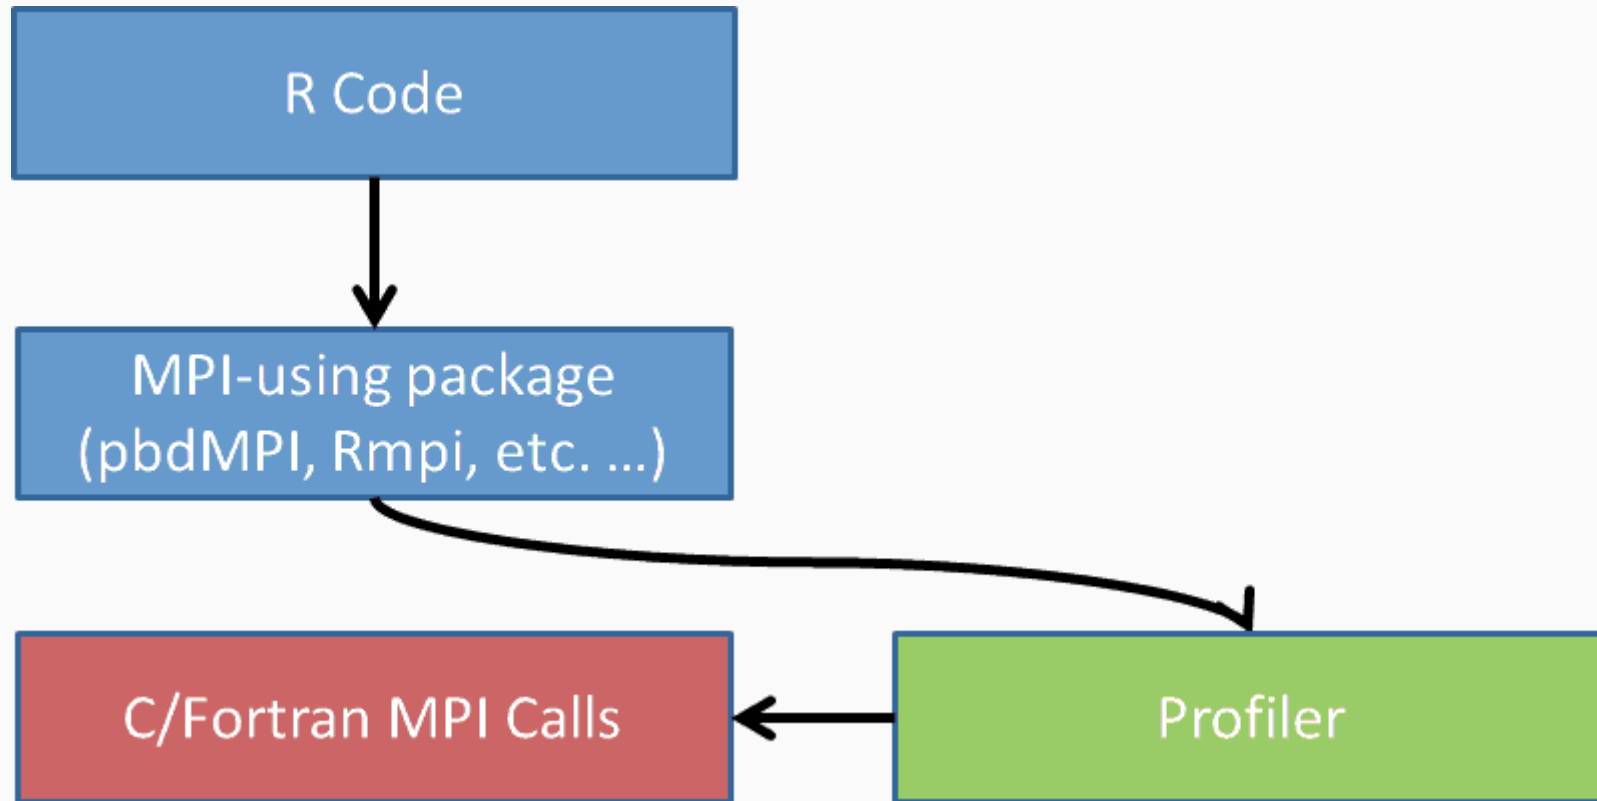


Hardware counters

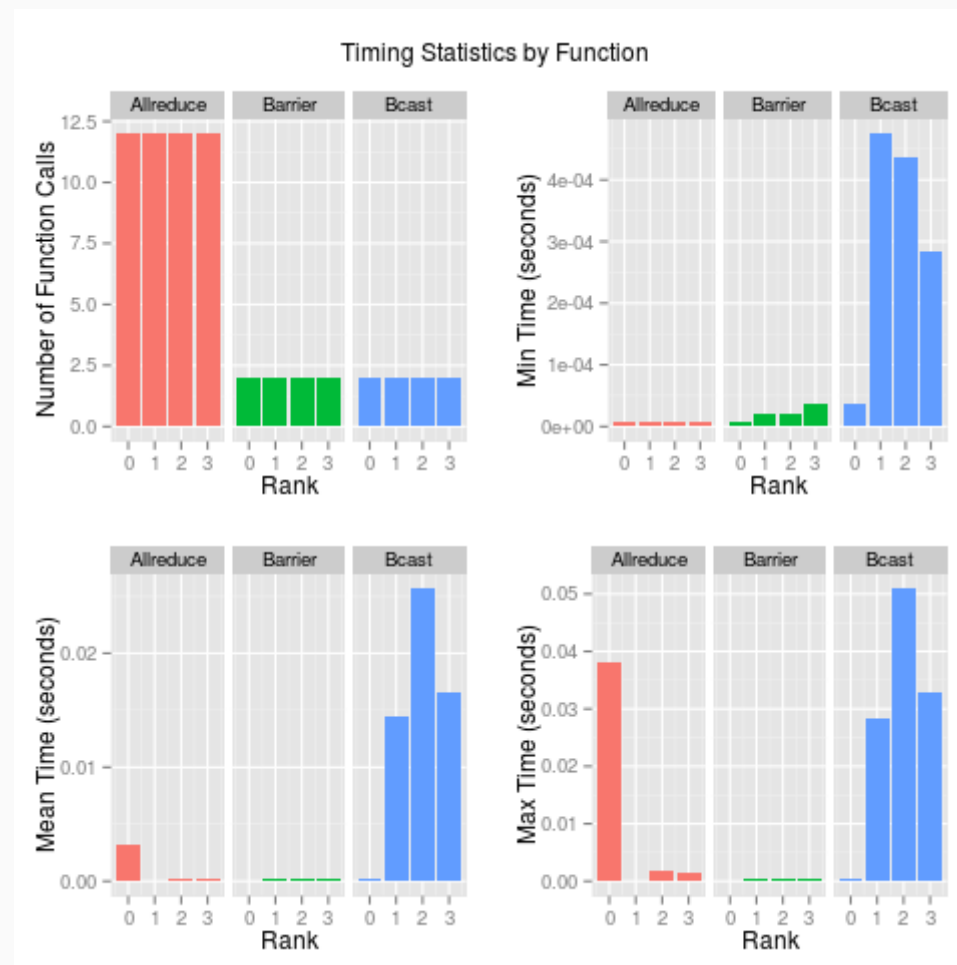
- Cache misses
 - Data cache
 - Instruction cache
- Flops
- Others



MPI Profiling



MPI Profiling



Questions?