

# Lecture 13 - Utilizing Compiled Code

DSE 512

---

Drew Schmidt  
2022-03-08

# From Last Time

- Next time: wrapup with I/O
- New homework "soon"
- Won't be due before March 26
- Questions?

# A Quick Comment


- I am not *ideological*
- But I am very *opinionated*
- I'm not always right
- You are free to disagree

# Compiled Code

- Very specialized set of skills
- Could *easily* spend *all semester* on just this
- This is a single lecture
- We will focus on just a few principles and examples
- This is probably the least useful topic for most of you
- *But it's important that you see it and understand it*

# Using Compiled Code in a HLL

- Pros
  - fast
  - memory-efficient
  - best of both worlds
- Cons
  - hard to write
  - hard to debug
  - multiple skillsets
- Julia???



# Differences

## HLL

- Memory managed *for* you
- Type (usually) implicit
- Interactive REPL

## LLL

- Memory managed *by* you
- Type (usually) explicit
- Non-interactive

# Memory

## Python

```
import numpy as np
np.array(range(0, n))
```

## C

```
double *x = malloc(n * sizeof(*x));
for (int i=0; i<n; i++)
    x[i] = i;
// ...
free(x);
```

# Types

R

```
x = 1  
typeof(x)
```

```
## [1] "double"
```

```
y = 2  
typeof(x)
```

```
## [1] "double"
```

```
z = 1:2  
typeof(z)
```

```
## [1] "integer"
```

C

```
double x = 1;  
double y = 2;  
int z = {1, 2};
```



# Bringing Compiled Code to HLL

- C/C++/Fortran/etc. code must be:
  - compiled
  - linked
  - loaded
- Each HLL handles this differently
- So why is this hard?

# Machine Code

```
#include <stdio.h>
```

```
int main()
{
    printf("hi\n");
    return 0;
}
```

```
gcc -g -O2 hello.c -o hello
```

(gdb) disass main

Dump of assembler code for function main:

```
0x00000000000001060 <+0>:    endbr64
0x00000000000001064 <+4>:    sub     $0x8,%rsp
0x00000000000001068 <+8>:    lea     0xf95(%rip),%rdi
0x0000000000000106f <+15>:   callq   0x1050 <puts@plt>
0x00000000000001074 <+20>:   xor     %eax,%eax
0x00000000000001076 <+22>:   add     $0x8,%rsp
0x0000000000000107a <+26>:   retq
```

End of assembler dump.

# Python Bytecode

```
def hello():  
    print('hi')  
  
import dis  
dis.dis(hello)
```

```
##      2          0 LOAD_GLOBAL          0 (print)  
##          2 LOAD_CONST          1 ('hi')  
##          4 CALL_FUNCTION        1  
##          6 POP_TOP  
##          8 LOAD_CONST          0 (None)  
##         10 RETURN_VALUE
```

# R Bytecode

```
hello = function() print("hi")  
hello = compiler::cmpfun(hello)  
compiler::disassemble(hello)
```

```
## list(.Code, list(12L, GETFUN.OP, 1L, PUSHCONSTARG.OP, 3L, CALL.OP,  
##      0L, RETURN.OP), list(print("hi"), print, structure(c(1L,  
## 9L, 1L, 30L, 9L, 30L, 1L, 1L), srcfile = <environment>, class = "srcref"),  
##      "hi", structure(c(NA, 0L, 0L, 0L, 0L, 0L, 0L, 0L), class = "expressionsIndex"),  
##      structure(c(NA, 2L, 2L, 2L, 2L, 2L, 2L, 2L), class = "srcrefsIndex")))
```

# Cython

I'm not going to be polite about this...



# How I Feel About Cython

- Only really useful for package developers
- Can't decide if it wants to be C or Python
- The Cython docs are *terrible*
- If you want a different take, see last year's Lecture 18

# Basic Structure

- `whatever.pyx` - Cython code
- `setup.py` - build script
- See the badly written docs for more information

[https://cython.readthedocs.io/en/latest/src/tutorial/cython\\_tutorial.html](https://cython.readthedocs.io/en/latest/src/tutorial/cython_tutorial.html)



# Example

There is no *basic* example I can give you that wouldn't be

- easier
- cleaner
- better

in Numba.

An Introduction to Just Enough Cython to be Useful

<https://www.peterbaumgartner.com/blog/intro-to-just-enough-cython-to-be-useful/>

# Rcpp

# Bringing Compiled Code to R

- R is a C program
- It has a C interface
- Do you know C?

- C++ interface
- Builds on top of the native C interface
- Can run compiled kernels "inline" (don't do this in production)
- Can be used in an R package (do this in production)

## What it is

- An R package
- A C++ interface for R extensions
- An R-like C++ interface
- One piece with many extensions
  - RcppArmadillo
  - RcppEigen
  - ...

## What it is not

- The only way to use C++ with R
- Zero-overhead
- As easy to use as R

# How I Think About Rcpp

Somewhere between Numba and Cython

*If you know what you're doing*

- You can often outperform Rcpp in runtime
- You can *always* outperform Rcpp in memory
- You can *massively* outperform Rcpp in compilation time

*and if you don't*

- Rcpp is probably your best bet



# Idiomatic C

```
void vec_vec_add(const int n, const double *a, const double *b, double *c){
    for (int i=0; i<n; i++)
        c[i] = a[i] + b[i];
}

void caller(const int n){
    double *a = malloc(n * sizeof(*a));
    double *b = malloc(n * sizeof(*b));
    set_values(n, a, b);

    double *c = malloc(n * sizeof(*c));
    vec_vec_add(n, a, b, c)
}
```

# R Native C Interface

```
#include <R.h>
#include <Rinternals.h>

SEXP vec_vec_add(SEXP a, SEXP b){
    R_xlen_t n = LENGTH(a);
    SEXP c;
    PROTECT(c = allocVector(REALSXP, n));
    for (int i=0; i<n; i++)
        REAL(c)[i] = REAL(a)[i] + REAL(b)[i];

    UNPROTECT(1);
    return c;
}
```

```
#include <Rcpp.h>

// [[Rcpp::export]]
Rcpp::NumericVector vec_vec_add(Rcpp::NumericVector a, Rcpp::NumericVector b){
    int n = a.length();
    Rcpp::NumericVector c(n);
    for (int i=0; i<n; i++)
        c[i] = a[i] + b[i];

    return c;
}
```

```
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
NumericVector vec_vec_add(NumericVector a, NumericVector b){
    int n = a.length();
    NumericVector c(n);
    for (int i=0; i<n; i++)
        c[i] = a[i] + b[i];


    return c;
}
```

# Monte Carlo $\pi$ Simulation

$$\text{Area} = \pi r^2$$

```
set.seed(1234)
n = 1000
x = runif(n)
y = runif(n)

circ_pts = seq(0, pi/2, length.out=100)
library(ggplot2)
g = ggplot(data.frame(x=x, y=y), aes(x, y)) +
  theme_bw() +
  geom_point() +
  annotate("path",
    x = 0 + 1*cos(circ_pts),
    y = 0 + 1*sin(circ_pts),
    color="red"
  )
```



# Estimating $\pi$ : Naive Implementation

```
pi_sim = function(n=1e6, seed=1234){  
  set.seed(seed)  
  
  s = 0  
  for (i in 1:n)  
  {  
    x = runif(1)  
    y = runif(1)  
    if (x*x + y*y < 1)  
      s = s + 1  
  }  
  
  4 * s / n  
}
```

```
system.time(pi_sim())
```

```
## user system elapsed  
## 3.883    0.012    3.898
```

# Estimating $\pi$ : Vectorized

```
pi_sim_vec = function(n=1e6, seed=1234){  
  set.seed(seed)  
  x = runif(n)  
  y = runif(n)  
  4 * sum(x*x + y*y < 1) / n  
}
```

```
system.time(pi_sim_vec())
```

```
##   user   system elapsed  
## 0.035    0.004    0.039
```

# Estimating $\pi$ : Rcpp

```
#include <Rcpp.h>

// [[Rcpp::export]]
double pi_sim_rcpp(int n){
  int s = 0;
  for (int i=0; i<n; i++){
    double x = R::runif(0, 1);
    double y = R::runif(0, 1);
    if (x*x + y*y < 1)
      s++;
  }

  return 4.0 * ((double) s / n);
}
```

```
Rcpp::sourceCpp("pi_sim_rcpp.cpp")
system.time(pi_sim_rcpp())[3]
```

```
##   user   system elapsed
## 0.019    0.000    0.018
```



# Summary

Method	Runtime	Relative Performance	Memory
Naive	3.898	216.556	~ 36 bytes
Vectorized	0.039	2.167	> 15 MB
Rcpp	0.018	1	36 bytes

# Questions?