

# Lecture 16 - Fork Parallelism Part 1

DSE 512

---

Drew Schmidt  
2022-03-24

# From Last Time

- Homework 3 posted tonight some time
- Questions?


# Parallelism Basics

# Vocabulary


- *Implicit parallelism* - Parallel details hidden from user
  - Example: Using multi-threaded BLAS
- *Explicit parallelism* - Some assembly required
  - Example: Using the `mclapply()` from the parallel package
- *Embarrassingly Parallel* aka *naturally parallel* aka *loosely coupled*  
Obvious how to make parallel; lots of independence in computations
  - Example: Fit two independent models
- *Tightly Coupled* - Opposite of embarrassingly parallel; lots of dependence in computations
  - Example: Speed up model fitting for one model

# Scalability

## Strong Scaling



## Weak Scaling



# Speedup

- *Wallclock Time* - Time of the clock on the wall from start to finish
- *Speedup* - unitless measure of improvement; more is better.

$$S_{n_1, n_2} = \frac{\text{Time for } n_1 \text{ cores}}{\text{Time for } n_2 \text{ cores}}$$

- $n_1$  is often taken to be 1
- In this case, comparing parallel algorithm to serial algorithm

# Independence

- Parallelism requires *independence*
- Ability to do different work at the same time
- Single problem split into many sub-problems
- Algorithm vs Implementation
  - Sometimes the algorithm is the hard part
  - Sometimes the parallel interface is
  - Sometimes it's both



# Independence Thought Experiment

- I have 10 laborers each with a shovel
- What if I need 10 holes dug?
- 5 holes?
- 20 holes?
- One really big hole?





# Independence in Practice

- R/Python function evals almost always independent
  - I/O may break this
  - Some HPC packages break this in memory (e.g. fmlr)
  - Other languages violate this all the time (e.g. C, Fortran, ...)
- Some examples by area:
  - Bio: split by gene
  - Chemistry/Materials: split by atom
  - Stats: split loglik sum
  - Bayesian: MCMC with multiple chains
  - ML: fit multiple models at once

# Overhead

- ***ALL*** parallel methods have overhead
- Sometimes it is undetectable
- Sometimes it is very noticeable
- Sometimes it is worse than serial!



# Fork

- Kernel-level parallelism mechanism
- Inexpensive process clone mechanism
- Only supported on Linux and MacOS
- Data is copy on modify!

# Fork in R and Python

- R: parallel package
- Python: multiprocessing package
- Both packages have other interfaces and backends (e.g. sockets)
- You can use that if you want (Windows), but we won't focus on it

# Fork Parallelism in R

`lapply()` becomes `mclapply()`:

```
mclapply(X, FUN, ..., mc.preschedule = TRUE, mc.set.seed = TRUE,  
         mc.silent = FALSE, mc.cores = getOption("mc.cores", 2L),  
         mc.cleanup = TRUE, mc.allow.recursive = TRUE, affinity.list = NULL)
```

# Basic Example in R

```
x = lapply(1:5, sqrt)
unlist(x)
```

```
## [1] 1.000000 1.414214 1.732051 2.000000 2.236068
```

```
ncores = parallel::detectCores()
x_mc = parallel::mclapply(1:5, sqrt, mc.cores=ncores)
all.equal(x, x_mc)
```

```
## [1] TRUE
```

# Fork Parallelism in Python

- A Bit more complicated...
- We'll try to keep it simple

# Basic Example in Python

```
import math
import multiprocessing as mp

list(map(math.sqrt, range(6)))
```

```
## [0.0, 1.0, 1.4142135623730951, 1.7320508075688772, 2.0, 2.23606797749979]
```

```
ncores = mp.cpu_count()
pool = mp.Pool(ncores)
pool.map(math.sqrt, range(6))
```

```
## [0.0, 1.0, 1.4142135623730951, 1.7320508075688772, 2.0, 2.23606797749979]
```



# "Portability"

- Not all OS's supported
- No fork on Windows
- Many API's try to hide this
- R and Python take different strategies
  - `mclapply()` errors; instead use `makeCluster()`
  - `Pool()` has a `Spawn()` method for Windows
  - ***DO NOT DESTROY YOUR WORKER POOL UNTIL YOU ARE DONE***




# Estimating $\pi$ in Python

# Monte Carlo $\pi$ Simulation

$$\text{Area} = \pi r^2$$

```
set.seed(1234)
n = 1000
x = runif(n)
y = runif(n)

circ_pts = seq(0, pi/2, length.out=100)
library(ggplot2)
g = ggplot(data.frame(x=x, y=y), aes(x, y)) +
  theme_bw() +
  geom_point() +
  annotate("path",
    x = 0 + 1*cos(circ_pts),
    y = 0 + 1*sin(circ_pts),
    color="red"
  )
```



# Serial Implementation

```
import random

def mc_pi(n):
    acc = 0
    for i in range(n):
        x = random.random()
        y = random.random()
        if x**2 + y**2 < 1:
            acc += 1
    return 4 * acc / n

random.seed(1234)
mc_pi(1000)
```

## 3.108

# Parallel Implementation

```
import multiprocessing as mp
import numpy as np
import random

def mc_pi_count(n):
    acc = 0
    for i in range(n):
        x = random.random()
        y = random.random()
        if x**2 + y**2 < 1:
            acc += 1
    return acc

def mc_pi_par(n, seed, ncores):
    random.seed(seed)
    pool = mp.Pool(ncores)
    acc = pool.map(mc_pi_count,
        np.repeat(round(n/ncores),
            ncores)
    )
    return sum(acc) * 4 / n
```

```
n = 10000000
for nc in [1, 2, 4]:
    t.start()
    pi_est = mc_pi_par(n, 1234, nc)
    t.stop()
```

```
## 4.66501244995743
## 2.4230600469745696
## 1.2336212289519608
```

# What About Numba?

```
from numba import jit

@jit(nopython = True)
def mc_pi_count(n):
    acc = 0
    for i in range(int(n)):
        x = random.random()
        y = random.random()
        if x**2 + y**2 < 1:
            acc += 1
    return acc
```

```
warmup = mc_pi_count(10)
n = 10000000
for nc in [1, 2, 4]:
    t.start()
    pi_est = mc_pi_par(n, 1234, nc)
    t.stop()
```

```
## 0.2649029800668359
## 0.09507097897585481
## 0.07486061903182417
```

# "Random" Numbers

- ~~RNG~~ PRNG
- Streams can overlap
- Be careful!
- Like all things:
  - Easy in R
  - Painful in Python




- *Random Number Generation* by Pierre L'Ecuyer  
<https://www.iro.umontreal.ca/~lecuyer/myftp/papers/handstat.pdf>
- R
  - It just works
  - Read `?parallel::mcparrallel` if you want
- Python
  - Depends *a lot* on the generator
  - Hard to do correctly
  - Better than nothing strategy: pass additional process-unique seed arg and use `starmap`
  - You're on your own



# Simulating Monty Hall in R

# The Monty Hall Problem

- You are shown 3 doors
- Behind 1 is a prize
- The other 2 have nothing/junk
- You choose a door
- One of the remaining 2 doors is shown to be junk
- You can switch to the remaining door if you choose
- What is the probability of winning if you switch?



# Simulating One Game

```
monty_hall = function(.)
{
  doors = 1:3
  prize_door = sample(doors, size=1)
  choice_door = sample(doors, size=1)
  # Always switch
  if (choice_door == prize_door)
    0 # lose
  else
    1 # win
}
```

```
set.seed(1234)
monty_hall()
```

```
## [1] 0
```

```
monty_hall()
```

```
## [1] 1
```

```
monty_hall()
```

```
## [1] 0
```

```
monty_hall()
```

```
## [1] 0
```

# Simulating Many Games

```
many_monty_hall = function(nruns)
{
  runtime = system.time({
    runs = lapply(1:nruns, monty_hall)
  })[3]

  c(mean(unlist(runs)), runtime)
}

set.seed(1234)
nruns = 1e6
many_monty_hall(nruns)
```


```
##           elapsed
## 0.667541 12.668000
```

# Going Parallel


```
library(parallel)
many_monty_hall = function(nruns, ncores)
{
  runtime = system.time({
    runs = mclapply(1:nruns, monty_hall, r
  })[3]

  c(mean(unlist(runs)), runtime)
}

set.seed(1234)
ncores = 1:12
nruns = 1e6
mmh = lapply(ncores, many_monty_hall, nrun
df = cbind(ncores, do.call(rbind, mmh)) |
as.data.frame()
```




# Parallel Speedup




# Going Parallel Again

```
set.seed(1234)
ncores = 1:12
nruns = 1e3
mmh = lapply(ncores, many_monty_hall, nruns)
df = cbind(ncores, do.call(rbind, mmh)) |>
  as.data.frame()
```



# Parallel Speedup Again





# Improving k-Means Clustering Outcomes in Parallel

# Parallelism in k-Means Clustering

- Parallelize the algorithm itself
  - Good for very large datasets
  - Tricky...
- Parallelize on the number of random starts
  - Easy to do!
  - May not be that valuable
- Parallelize on the number of clusters
  - Also easy to do!
  - Turn your brain off! Machine your learnings!!!

# Measuring Cluster Performance

- Would like an analogue of classification accuracy
- How do we compare two sets of cluster labels?
- Rand Measure / Rand Index [https://en.wikipedia.org/wiki/Rand\\_index](https://en.wikipedia.org/wiki/Rand_index)
  - Given a set of  $n$  elements  $s$  and two partitions of  $s$  into  $r$ -length subsets  $X = \{X_1, \dots, X_r\}$  and  $Y = \{Y_1, \dots, Y_r\}$

$$R = \frac{a + b}{\binom{n}{2}}$$

Where


- $a$  is the # of pairs of elements of  $s$  that agree across  $X$  and  $Y$
- $b$  is the # of pairs of elements of  $s$  that disagree across  $X$  and  $Y$

# Rand Measure

```
rand_measure = function(l1, l2){  
  n = length(l1)  
  a = b = 0L  
  
  for (i in 1:n){  
    for (j in (i+1L):n){  
      if (j > n) # R indexing is stupid  
        break  
  
      same1 = (l1[i] == l1[j])  
      same2 = (l2[i] == l2[j])  
  
      if (same1 && same2)  
        a = a + 1L  
      else if (!same1 && !same2)  
        b = b + 1L  
    }  
  }  
  
  (a + b) / choose(n, 2)  
}
```

# Using the Iris Dataset


- Fisher's famous iris dataset
- Required by law in clustering lectures



# Finding the k in k-Means

```
library(parallel)
labels = as.numeric(iris[, 5])
x = iris[, -5]

options(mc.cores=8)
nclusters = 2:100
clusters = mclapply(
  nclusters,
  function(i) kmeans(x, i, nstart=10)$clusters
)
rand = mclapply(
  clusters,
  rand_measure,
  l2 = labels
) |> unlist()
```



# Wrapup

# Scalability?

- "scale up"
- "scale out"
- What do these mysterious phrases mean?





# Vertical (up) vs Horizontal (out) Scaling

- Industry jargon
- (Opinion) Not really that well-defined
- More about *hardware* than *applications*
- Example
  - *Scale up* - Get a better processor
  - *Scale out* - Get multiple nodes
- Have we seen either so far?

# Wrapup

- Be careful with RNG's in parallel!
- Where to look for (task) parallelism
  - Independent simulations
  - Multiple model fitting
  - *Many* ETL operations
- "Workflow parallelism" is *ubiquitous* in data/analysis pipelines
- Next time: data parallelism

# Questions?