

Lecture 25 - Advanced Profiling

DSE 512

Drew Schmidt
2022-04-28

From Last Time

- Homework 4
 - Assigned
 - Due April 30
 - Extensions very unlikely
- Homework 5
 - Covers profiling
 - Smaller
 - Assigned next week?
 - Due May 10?
- No Homework 6 (Deep Learning)
- Questions?

Advanced Profiling

Advanced Profilers

- Hardware counter profiling
- MPI profilers
- GPU (CUDA) profilers

Hardware Counter Profiling

- Modern hardware (CPU, GPU) collects *hardware counters*
- Counter: number of occurrences
- Hardware counter: Number of times hardware event occurs
- Performance Application Programming Interface (PAPI)

<https://icl.utk.edu/papi/>



Hardware Counter Examples

- Cache misses
 - Data cache
 - Instruction cache
- Flops
- Others



An Example

We're talking about SVD again!



- Input matrix A
 - Compute SVD of A
 - Project A onto V
- Number of floating point operations:
 - SVD requires $6mn^2 + 20n^3$ ops
 - projection requires $2mn^2$ ops
 - Source: Golub, G.H. and Van Loan, C.F., 2013. Matrix computations. JHU press.
 - Also requires mean-centering: $2mn + n$ ops
- TOTAL: $6mn^2 + 20n^3 + 2mn^2 + 2mn + n$

PCA FLOPs

```
m = 10000
n = 50
A = matrix(rnorm(m*n), m, n)
ret = prcomp(A)

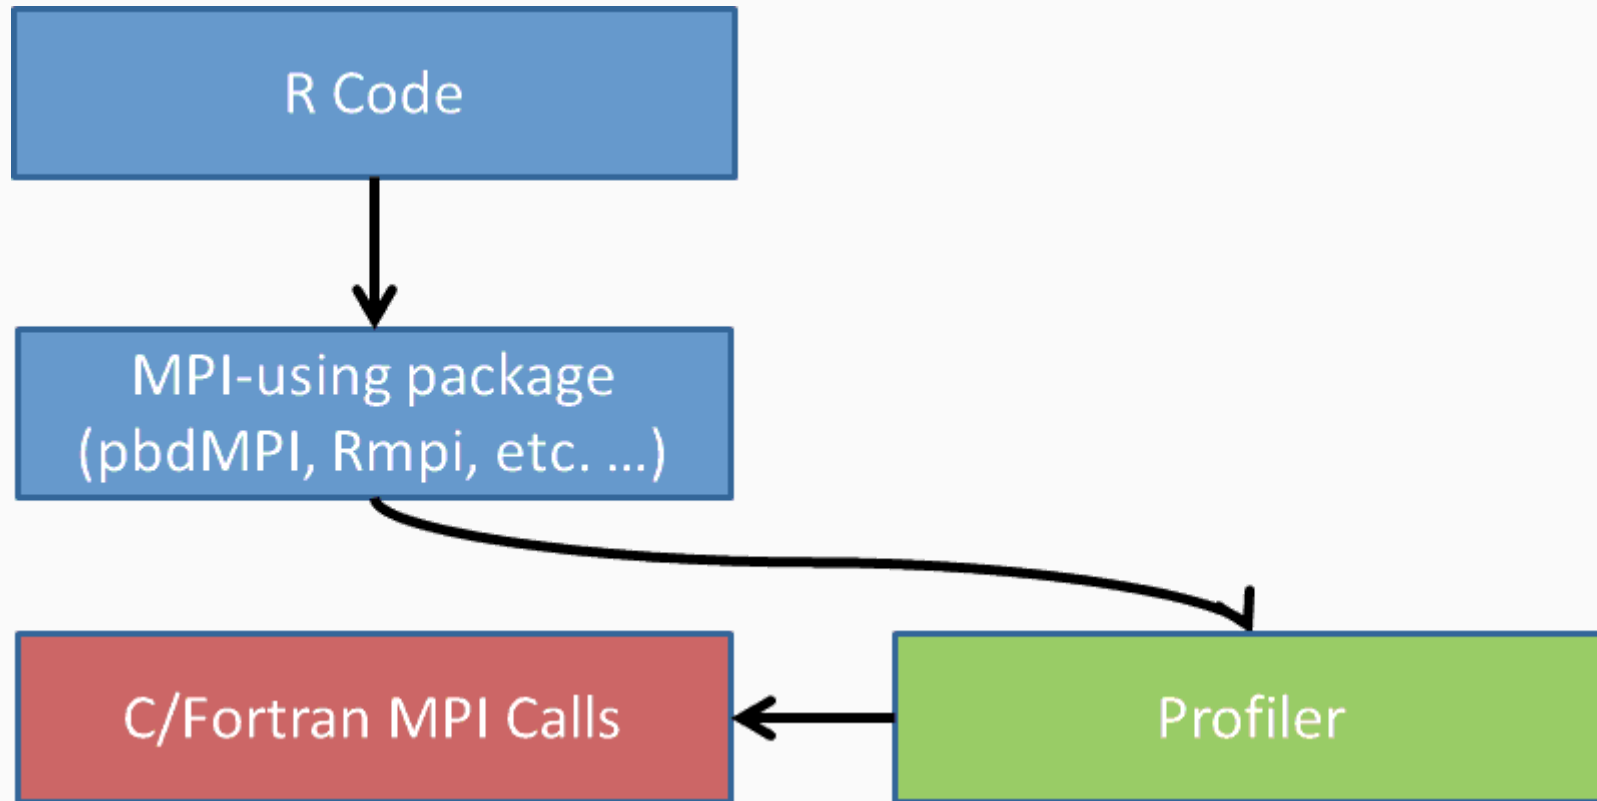
pbdPAPI::system.flops(prcomp(A))
```

	m	n	measured	theoretical	difference	pct.error	mflops
1	10000	50	212538720	203500050	9038670	4.25	2284.26

MPI Profilers

- fpmapi
 - Easy to install
 - Least useful
- mpiP
 - Relatively easy to install
 -
- tau
 - *Very* hard to install
 - Profiles *everything*

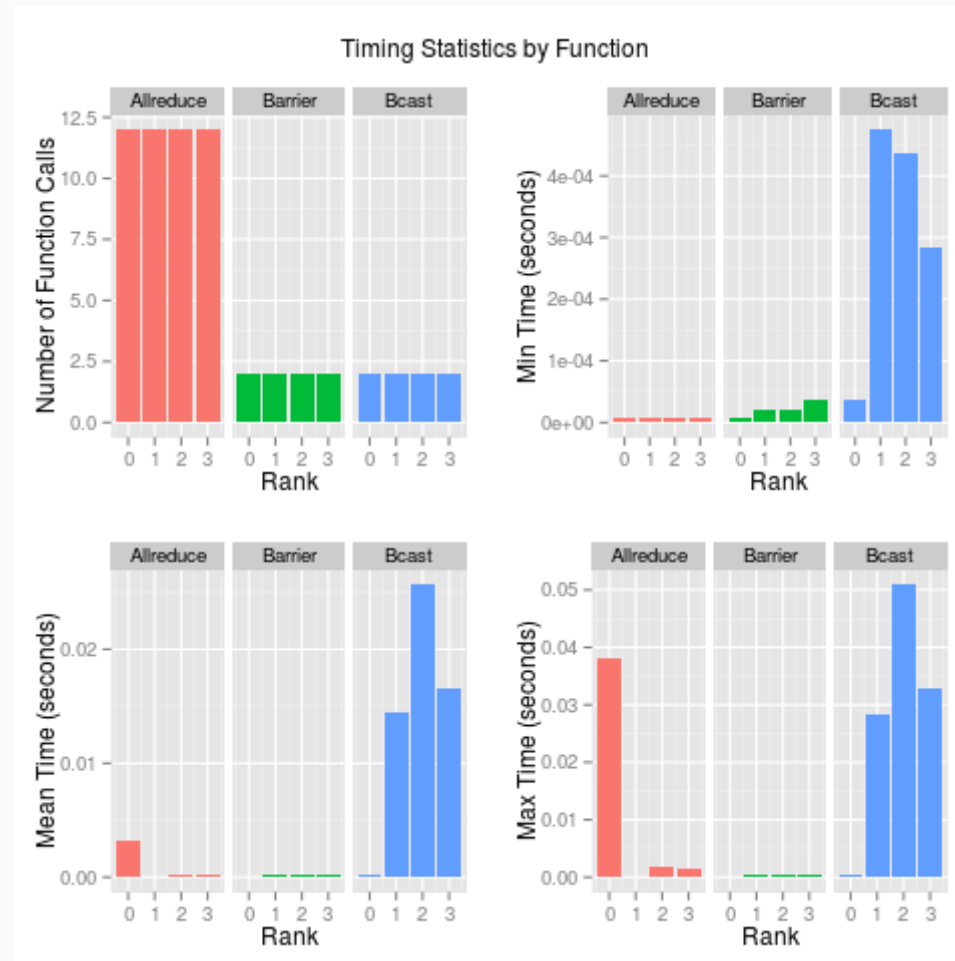
MPI Profiling



MPI Profiling Steps

- Build mpiP
- Rebuild your MPI program
 - Binary, pbdMPI, mpi4py, ...
 - Link with mpiP
- Run program
- Read profiler data

MPI Profiling



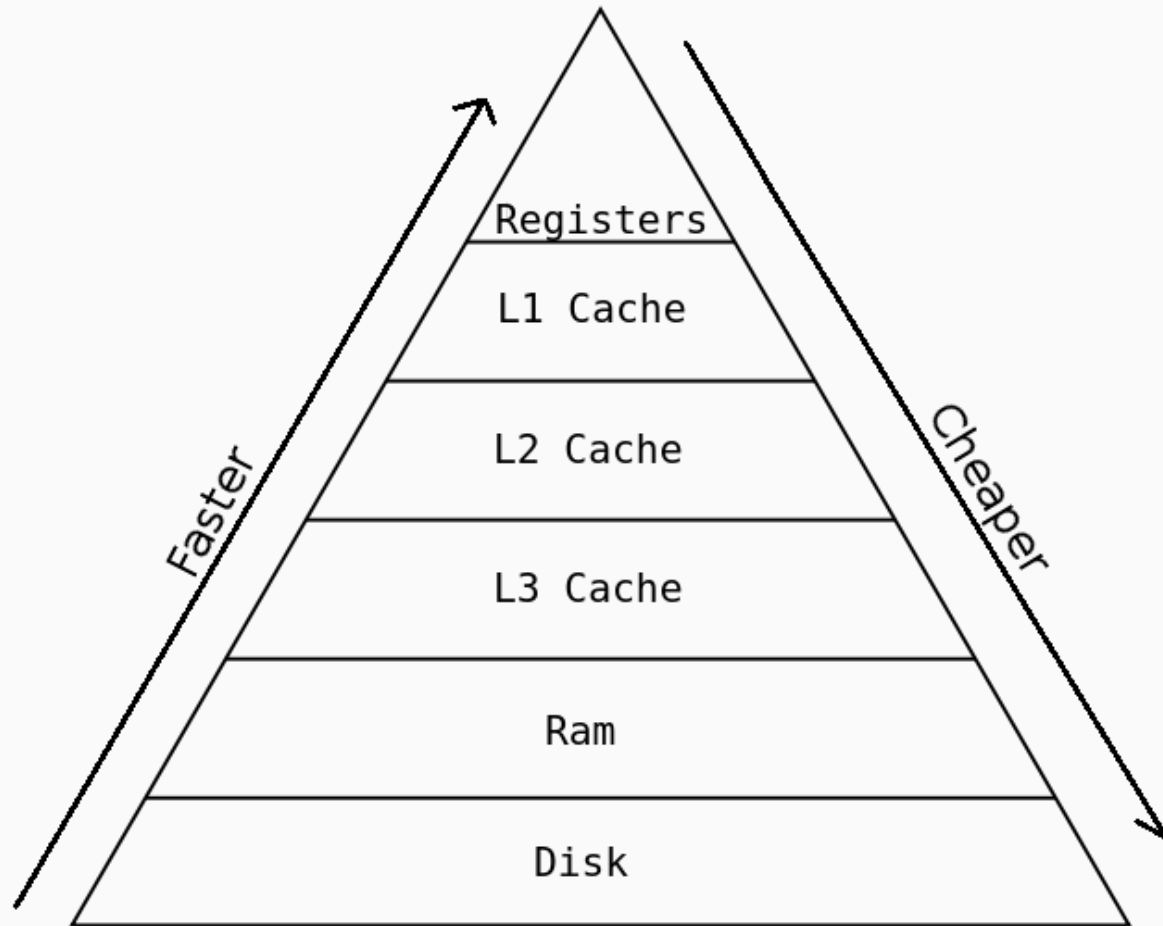
Cache

Computer Hardware

- We have to talk a bit about computer hardware
- Goal is not to become experts
- Need to be aware of some realities

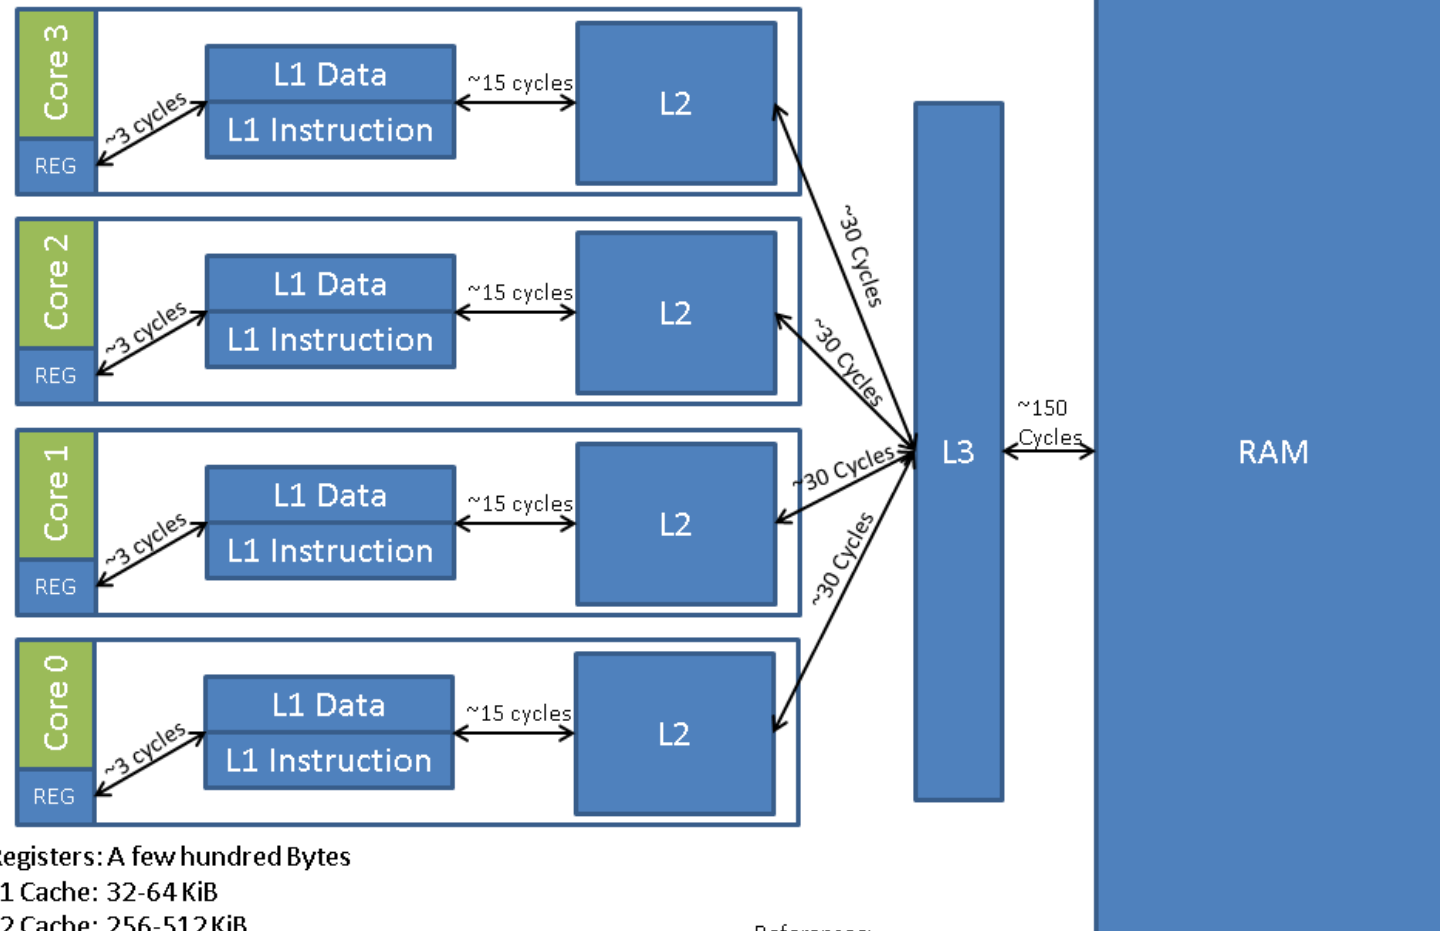


Computer Memory Hierarchy



Cache Hierarchy*

*sizes not drawn to scale



Registers: A few hundred Bytes

L1 Cache: 32-64 KiB

L2 Cache: 256-512 KiB

L3 Cache: 4-8 MiB

Ram: 8-16 GiB

References:

http://www.eecs.berkeley.edu/~rcs/research/interactive_latency.html

<http://www.7-cpu.com/cpu/SandyBridge.html>

Cache Sizes

```
library(memuse)
```

```
Sys.cachesize()
```

```
## L1I:    64.000 KiB
```

```
## L1D:    32.000 KiB
```

```
## L2:    512.000 KiB
```

```
## L3:     16.000 MiB
```

```
Sys.cachelinesize()
```

```
## Linesize:  64 B
```

Recall: Homework 3

Filling matrix $A = [a_{ij}]_{n \times n}$ where

$$a_{ij} = \exp\left(-\frac{i+j}{ij}\right)$$

Homework 3

R

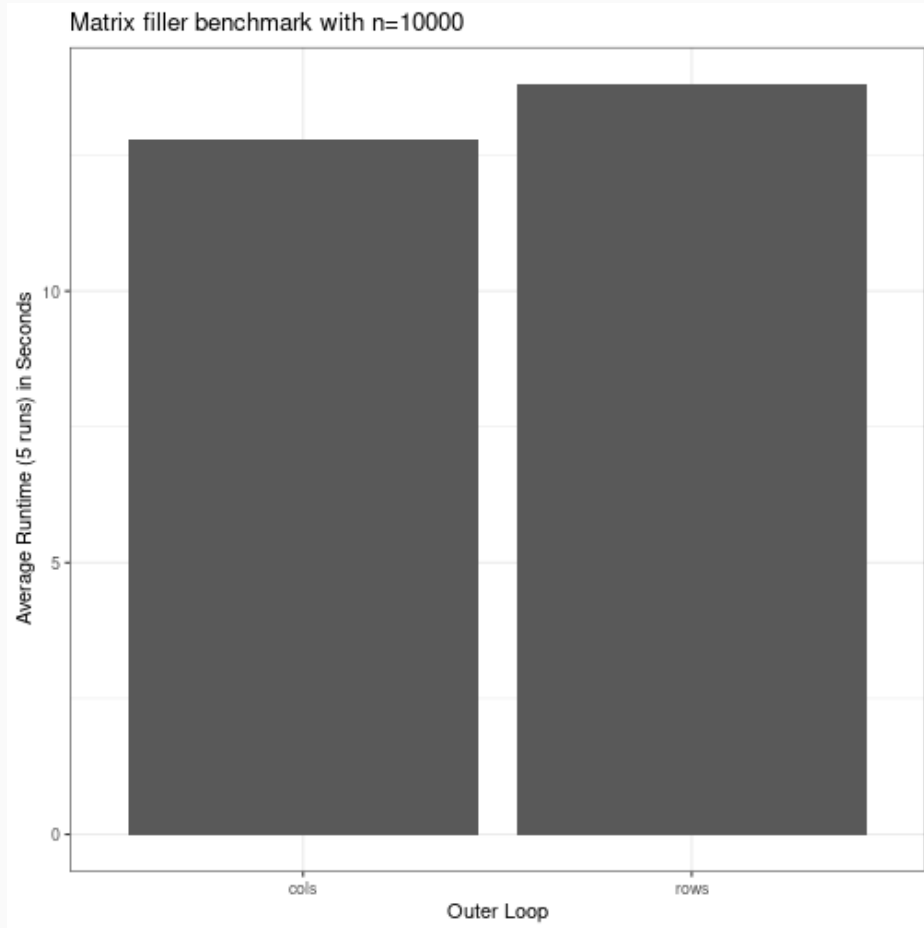
```
rf = function(n){  
  A = matrix(0, n, n)  
  for (i in 1:n)  
    for (j in 1:n){  
      A[i, j] = exp(-(i+j)/(i*j))  
    }  
  A  
}  
  
cf = function(n){  
  A = matrix(0, n, n)  
  for (j in 1:n){  
    for (i in 1:n)  
      A[i, j] = exp(-(i+j)/(i*j))  
    }  
  A  
}
```

Python

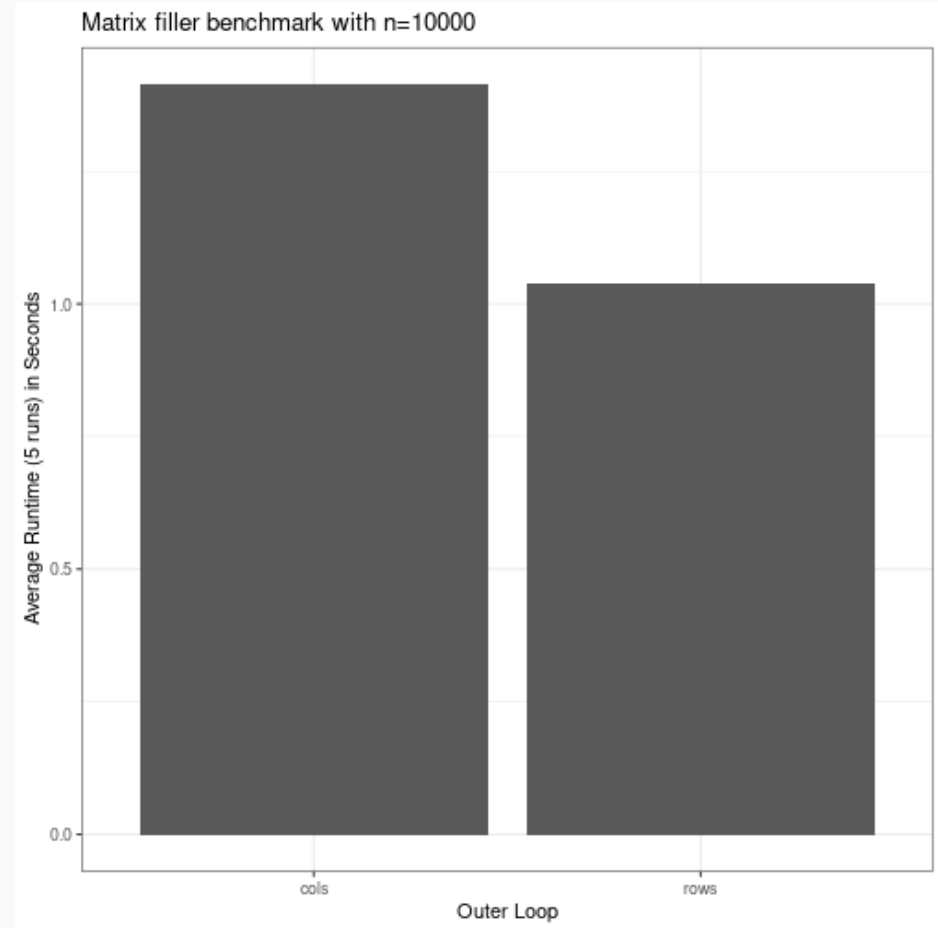
```
@jit(nopython = True)  
def rf(n):  
    A = np.zeros((n, n))  
    for i in range(0, n):  
        for j in range(0, n):  
            A[i, j] = math.exp(-((i+1)+(j+1))/(  
        return A  
  
@jit(nopython = True)  
def cf(n):  
    A = np.zeros((n, n))  
    for j in range(0, n):  
        for i in range(0, n):  
            A[i, j] = math.exp(-((i+1)+(j+1))/(  
    return A
```

Homework 3

R



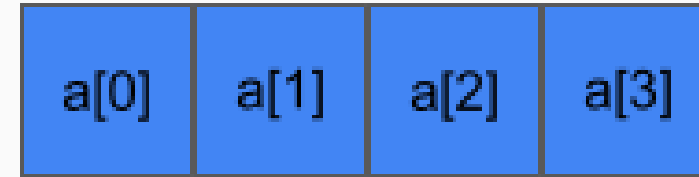
Python (with Numba)



Why Would It Matter?

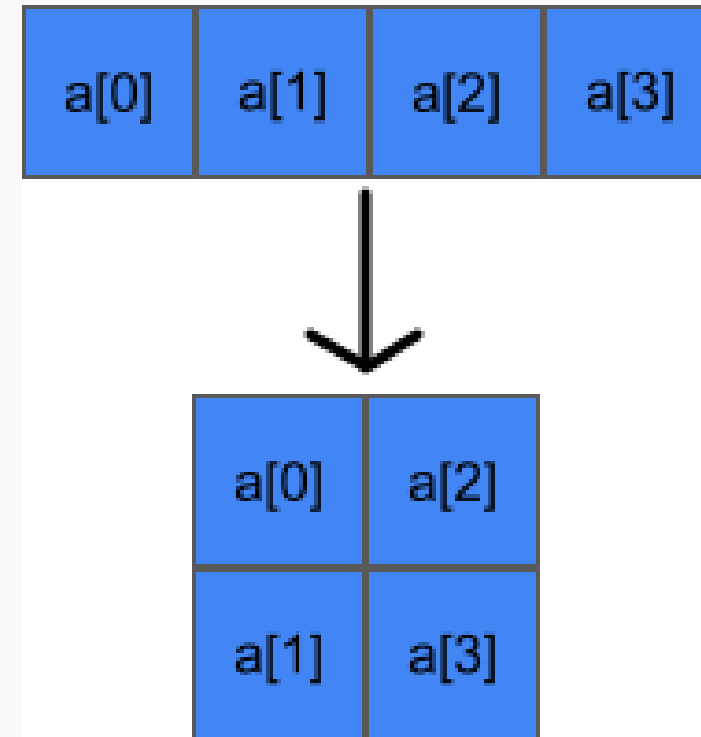
Matrices

- Computers don't have a concept of "matrix"
- Computers *do* have a concept of "array"
- Array: A contiguous block of memory



Matrices

- Matrices are (almost always) just arrays
 - Can be arrays of arrays
 - This is usually not a good idea
 - I don't want to get into it
- Ordering (column-major / row-major)
- Col: $A[i, j] = A[i + m*j]$
- Row: $A[i, j] = A[i*n + j]$



Looping "Wrong"

- CPUs will fetch data to cache(s) that it expects you'll need
- A cache "miss" means data is not in cache
 - have to go get it from RAM
 - RAM → L3 → L2 → L1
- A "miss" is not inherently bad
 - If everything fit in cache, GREAT!
 - Probably it won't
 - In that case, *there will be cache misses*
- What you *don't* want is **unnecessary** misses

Another Example

We're simulating π again!



Implementations

Rcpp

```
double mcpi_rcpp(const int n){
    int i, r = 0;
    double u, v;
    Rcpp::RNGScope scope;

    for (i=0; i<n; i++){
        u = R::runif(0, 1);
        v = R::runif(0, 1);
        if (u*u + v*v <= 1)
            r++;
    }

    return (double) 4.0*r/n;
}
```

C

```
SEXP mcpi_c(SEXP n_){
    SEXP ret;
    int i, r = 0;
    const int n = INTEGER(n_)[0];
    double u, v;

    GetRNGstate();
    for (i=0; i<n; i++){
        u = unif_rand();
        v = unif_rand();
        if (u*u + v*v <= 1)
            r++;
    }
    PutRNGstate();

    PROTECT(ret = allocVector(REALSXP, 1));
    REAL(ret)[0] = (double) 4.0*r/n;
    UNPROTECT(1);
}
```

Benchmark

```
library(Rcpp)
```

```
sourceCpp(code=Rcpp_code)
```

```
sourceCpp(code=C_code)
```

```
n = 1e6
```

```
rbenchmark::benchmark(mcpic(n), mcpircpp(n))
```

##	test	replications	elapsed	relative	user.self	sys.self	user.child
## 1	mcpic(n)	100	1.609	1.000	1.598	0.004	0
## 2	mcpircpp(n)	100	3.099	1.926	3.084	0.001	0

Data Cache

```
pbdPAPI::system.cache(mcpic(n))
```

```
## Level 1 Cache Misses: 1358  
## Level 2 Cache Misses: 720  
## Level 3 Cache Misses: 400
```

```
pbdPAPI::system.cache(mcpircpp(n))
```

```
## Level 1 Cache Misses: 1316  
## Level 2 Cache Misses: 747  
## Level 3 Cache Misses: 248
```

Instructions

Total instructions executed

```
pbdPAPI::system.event(mcpic(n), events="PAPI_TOT_INS")
```

```
## Instructions Completed: 131035426
```

```
pbdPAPI::system.event(mcpircpp(n), events="PAPI_TOT_INS")
```

```
## Instructions Completed: 227028352
```

Instruction Cache

Instruction cache misses

```
pbdPAPI::system.event(mcpic(n), events="PAPI_L1_ICM")
```

```
## Level 1 Instruction Cache Misses: 604
```

```
pbdPAPI::system.event(mcpircpp(n), events="PAPI_L1_ICM")
```

```
## Level 1 Instruction Cache Misses: 873
```

Wrapup

Wrapup

- Profiling is information gathering
 - runtimes
 - memory consumption
 - cache misses
 - MPI comms
- The more you understand your computer's memory hierarchy, the faster your programs will be.
- That's it for profiling!
- Deep learning next

Questions?